





Recent AI model trade-offs: intelligence, throughput, and price

A. Three-objective Pareto overview



- Pareto candidates (IDs used in panels B-D); ★ marks the balanced candidate
- |                     |                         |                               |
|---------------------|-------------------------|-------------------------------|
| ● M02 GPT-5.6 Terra | ● M06 Grok 4.6          | ★ M13 Gemini 3.7 Flash        |
| ● M03 GPT-5.6 Luna  | ● M07 Qwen3.8 Max       | ● M15 Nemotron 3.5 Lightning  |
| ● M04 Claude Opus 5 | ● M08 DeepSeek V4 Flash | ● Dominated model (unlabeled) |

Source: Artificial Analysis, retrieved 2026-08-14. 16 models; 8 nondominated. Pareto status uses raw measured values. Price axes use a logarithmic scale; arrows show the preferred direction. Blended price is Artificial Analysis's 7:2:1 cache-hit/input/output token mix.